

A business case for Concord

I'm Xavier — a software engineer who likes **solving hard problems** and **owning them end to end**.

HOW I TAKE OWNERSHIP — END TO END




That's the builder in me — but most of my time here went into understanding **your** world, not the code. So here's my map; please correct me.

Xavier
Medina · Software engineer — 4+ yrs shipping
production AI

· ProsperaLabs — voice · chat · email · ecommerce, real-estate, legal &
customer-service

Four strategies, one focus — Concord starts with your direct PE portfolio companies




v1 = PE only; the other three branches are **documented extensions, not built.**

Re-verify Sagard's strategy line-up + scale figures on sagard.com the morning of; print no AUM / portfolio-count / fund-size numbers as fact.

2 · Concord · prepared for Sagard

Fundraise → invest → hold → exit: Concord lives in the quarter-by-quarter middle



A fund lives ~10 years; monitoring is the quarter-by-quarter middle.

Concord feeds the inputs to **the mark** — it never makes the mark (a human, audited judgment).

Five people read the same pack — for opposite jobs

	PERSONA	WHAT THEY DO	TRUST POSTURE
1	Deal partner	Fast “on track?” board read	Speed over precision
2	Reporting Operations + Controls	Owens the reconciliation & reporting grind	Comparability + control ← our customer 
3	Valuation analyst	Sets each mark → fund NAV	Precision above all
4	IR / LP-reporting	Reports numbers outward	Precision (external)
5	4 · Concept prepared for Sargent Risk / Credit analyst	Checks borrowers stay in covenants	Precision + exact definitions

Same builder, opposite tools — speed for the deal partner, precision for Ops

ROUND 1 · LLM · SPEED

Deal partner → “AI Scout”

- company URL → scrape info + press
- search + scrape the founders
- summarize → Slack: “move forward?”

the deal partner wants **speed** — a fast, directional read that keeps a deal moving

🕒 Concord

ROUND 2 · DETERMINISTIC · PRECISION


Operations → Concord

- a number you'd sign your name to
- traceable to its source
- a wrong answer refused, not guessed

precision first — that's why this one is deterministic

Same builder, two rounds, **two opposite trust postures**.


Monitoring the portfolio means hand-keying 20+ packs — every quarter



Synthetic test packs — 10 companies across 24 PDFs. “20+ packs,” not “20+ companies.”

6 · Concord · prepared for Sagard

How Concord works today, end to end



Reading the number is the easy **10%**; everything below — comparability + trust — is the expensive **90%**.

This is what runs **today** — one deterministic pass; provenance is **file-level**. The redundant / orchestrator version is the future (deck slide 17).

Four traps in the manual close — each one provable in the packs

1 · SAME METRIC, MANY NAMES

NovaCloud ARR appears as “End-of-Period ARR” and “ARR (End of Period)” — two labels, one line. Miss a relabel and the trend silently splits.

2 · SAME NAME, DIFFERENT MEANING

SaaS delivery-margin vs **LendBridge** interest-margin — put them in ~~one shared column~~ and you're quietly wrong.

3 · RIGHT NUMBER, WRONG BASIS

ClearPay “Cash \$38.4M” includes \$6.2M restricted client float — true spendable is ~~\$38.4M~~ \$32.2M.

4 · CAN'T PROVE IT

The sheet is stale, retyped — **no number links back to its source.** “Where did this come from?” → “let me re-open the PDF.”

Same label ≠ same metric — comparability is the product, not extraction

SAAS COMPANY

60%

“gross margin” (illustrative) — a delivery margin

=?

LENDER

60%

“gross margin” (illustrative) — an interest margin

THE NAIVE MERGE ~~one shared “gross margin” column~~ splits →

SAAS LANE delivery-margin | LENDER LANE interest-margin

THE INVERSE one metric, many names — “End-of-Period ARR” = “ARR (End of Period)” →

one stitched line

Illustrative numbers. Live, this fires on **LendBridge's** real interest-margin (a lender's “gross margin,” 58 → 62%) — shown in full, never ranked against a SaaS margin.

Rules never invent a number — deterministic on purpose, not production-grade yet

ON PURPOSE

A single deterministic method — fixed rules

it never invents a number — a flagged blank beats a confident guess

cheap

fast

testable

offline

THE HONEST CAVEAT

One method ≠ production-grade


a lone reader, however good, can be wrong — and you'd never know.
The redundant path (§8) is what I'd build for real.

FIRECRAWL → DETERMINISTIC

first build used Firecrawl (paid extraction) → hit errors + ran out of credit → switched. The right move for a **trust** tool.

Any vendor accuracy % = the vendor's own claim, not our measured result.

What I'd build for real: redundant readers — the disagreement is the alarm



Designed, not built — today's demo runs the single deterministic method. “Independent” is never total: agreement is strong evidence, not proof.

Be AI-first, but human-decision-driven — AI never decides

AI / RULES · DOES THE HEAVY WORK

- read 20+ packs, pull the numbers
 - stitch renamed labels into one series
 - flag disagreement & low confidence
 - suggest — **with reasons**
- probabilistic — a likely answer, not a certain one



SUGGESTS

THE HUMAN DECISION GATE · DECIDES

- “is this right enough for a report?”
 - sign-off
 - the mark / the IC call
- the center of every decision that reaches a report

The win isn't the hours — it's a wrong number caught before the IC

TODAY

Hours-days

re-typing 20+ packs by hand before anyone can even see the portfolio

WITH CONCORD

One click

a comparable, source-traced view — time goes to judgment, not typing

typing saved

a wrong number caught before the investment committee — worth far more

Qualitative only — no invented hours, no dollar figure, no fund-return math.

► **LIVE DEMO**



One comparable, source-traced view of every
portfolio company — run live and offline.

offline

deterministic

provenance-tagged

Run live from the app · the three beats and every caveat are spoken, not on slides

Reliability isn't more AI — it's more agreement, more human judgment, and a record of both



More agreement

independent readers that fail differently — the ensemble raises warnings



More human judgment

the machine flags; a human decides every call that reaches a report




A record of both

a permanent, exportable trail — who decided what, and why

Everything in this section is **designed, not built** — four I'll walk through, all roadmap.


Correlated readers can outvote the one that's right — so independence, not count, decides



3 copies of one model share blind spots — **confidently wrong the same way** (correlated failure). Independence, not count, decides.

Designed, not built — no auto-accept rate until it runs on real packs. Reducto ~99% and any vendor figure is the vendor's own claim.

The architecture in one picture



All three readers are AI — agreement is **strong evidence, not proof**, so the disagree path ends at a **person**. Every fix is logged and teaches the readers. **Future expansion** — designed, not built.

AI writes the interpretation, never the number — and a human can always refute it

1 · CLEAN GRID

already-normalized, provenance-tagged numbers

→

AI SUGGESTION — REVIEW

NovaCloud: compounding +42% ARR, but cash falling \$29.5M → \$19.6M — confirm the funding plan.

→

3 · HUMAN

✓ accept ✎ correct 💬 comment

🔄 **Loop:** next quarter the AI reads last quarter's comments first — it learns the house view.

Roadmap feature — the AI note is not built; it can only get an interpretation wrong (the human refutes it), never inject a wrong number. Cites every number · no dollar-ROI claims.

Concord refuses to mix pounds and dollars — it flags, it never converts

TODAY

A front-end rule flags GBP

PeopleFlow: “· not comparable (GBP)” — kept off the dollar axis. The parser sees **currency = null**; a human-owned rule catches the £.


ROADMAP

A currency-reader captures the £

- convert **levels only** — dated rate (fx_rate · fx_source · fx_date)
- refuse to convert ratios / % — a GBP % stays a %
- refuse over convert where the basis is unknown

The GBP flag is a front-end rulebook rule, **not parsed data** (parser sees currency = null). No FX conversion exists yet.

No silent overwrites, maker-checker on every fix, an audit trail that stands up to an LP challenge



None of this exists yet — the next layer on top of provenance + refuse-to-compare, which **do** exist.

Concord today is a prototype, not a production system

WHAT'S REAL

- the offline, deterministic engine
- recall **76% → 90%, 0 wrong values** (24 synthetic PDFs)
- refuse-to-compare across asset classes
- file-level provenance

NOT YET — TO ADD

- a durable store
- a job queue + retries
- the redundant ensemble
- multi-tenant cloud

prototype, not production


The 24 PDFs are synthetic test fixtures, not real filings. The redundant ensemble is **designed, not built** — “to add,” never present-tense.

Harden in place first — the cloud is the destination, when it's justified

	OPTION 1 — HARDEN IN PLACE	OPTION 2 — CLOUD RE-PLATFORM
What it is	Keep the engine; add a durable store + job queue/retries + the redundant ensemble; stay offline/on-prem	Object storage → file event → queue fans out → managed workers → DB → publish; auto-scales for bursts
Reliability	High gain fast — the ensemble is the biggest single lever	Highest ceiling — but only after real hardening + monitoring
Risk	Low — additive + reversible; deterministic core stays a canary	Med-High — data leaves on-prem (a data-residency question)

Do Option 1 first; Option 2 is the destination. The ensemble inside Option 1 is designed, not built — “to add.” Depth in iv-system-design-and-scaling.md.

Most of my time wasn't coding — it was understanding your world



Xavier's own rough estimate — self-reported, not measured. That's the FDE job: point the code at the **right** problem.

Comparability, provenance, a path forward — now, your questions



Comparability is the product

not extraction — knowing when not to put
two numbers in one column



Provenance you can prove

every number one click from its source



A clear path forward

to the reliable, redundant, production-
grade version

Over to you — and I have a few questions for you.

PREPARED FOR



One comparable, source-traced view of
every portfolio company.

Built with AI, decided by humans ❤️

Made with love, with AI — but human-first on every decision. Thanks.

Xavier Medina · Forward Deployed Engineer · xmedinavei@gmail.com